

Agentic AI Frameworks for Autonomous Cloud Resource Management

Folakemi Idowu¹ | Emmanuel Idowu²

¹ Department of Computer Science, Lagos State University (LASU)

² Department of Computer Science, Ladoke Akintola University of Technology (LAUTECH)

folakemi.idowu230910077@st.lasu.edu.ng¹ | emmybabm@gmail.com²

Date: 5 March 2026

ABSTRACT

Cloud computing infrastructure has grown in complexity to a degree that manual administration and rule-based automation are no longer adequate to ensure optimal performance, cost efficiency, and fault tolerance. This paper examines agentic artificial intelligence (AI) frameworks as a paradigm for fully autonomous cloud resource management. Agentic AI systems combine large language model (LLM) reasoning, multi-agent orchestration, reinforcement learning, and policy-driven execution to enable cloud environments to perceive their operational state, formulate management strategies, and enact resource allocation decisions without continuous human supervision. We survey the architectural constituents of agentic cloud systems — including perception modules, reasoning engines, action executors, memory subsystems, and inter-agent communication protocols — and assess how these components interact within modern cloud-native stacks. Drawing on current literature and empirical evidence from production deployments, we analyze the capabilities of agentic frameworks across five functional domains: dynamic auto-scaling, cost optimization, anomaly detection and self-healing, workload scheduling, and compliance assurance. We further identify the key challenges associated with deploying agentic AI in cloud environments — notably safety constraints, explainability, adversarial robustness, and governance — and propose a research agenda for addressing these gaps. Our analysis demonstrates that agentic AI frameworks represent a transformative advancement in cloud resource management, capable of delivering substantial reductions in operational overhead, infrastructure cost, and mean time to recovery (MTTR) compared with conventional automation approaches.

Keywords: Agentic AI, autonomous cloud management, multi-agent systems, LLM-based orchestration, reinforcement learning, cloud-native architecture, DevOps automation, self-healing systems, cost optimization, AI governance.

1. INTRODUCTION

The rapid proliferation of cloud-native applications, microservices architectures, and containerized workloads has introduced an unprecedented degree of operational complexity into modern IT infrastructure. Enterprises now routinely manage thousands of interdependent services across hybrid and multi-cloud environments, where resource demands fluctuate dynamically in response to user traffic, batch processing schedules, and external events. Traditional approaches

to cloud resource management — ranging from manually configured autoscaling rules to threshold-based alerting systems — were designed for environments of substantially lower complexity and are increasingly inadequate to meet the reliability, efficiency, and agility requirements of contemporary cloud deployments.

Automation frameworks such as Kubernetes Horizontal Pod Autoscaler (HPA), AWS Auto Scaling Groups, and Terraform-managed infrastructure-as-code have partially addressed this challenge by codifying operational expertise into reusable policies. However, these systems operate within fixed parametric boundaries and lack the capacity for contextual reasoning, cross-domain synthesis, or adaptive strategy formation in response to novel operational conditions. The result is a persistent gap between what automated systems can handle predictably and what production environments actually demand.

Agentic AI represents a fundamentally different approach to this problem. Rather than encoding operational knowledge as static rules or preconfigured policies, agentic systems leverage the generative reasoning capabilities of large language models (LLMs), the strategic optimization power of reinforcement learning agents, and the collaborative capacity of multi-agent orchestration frameworks to achieve goal-directed behaviour in dynamic, partially observable environments. An agentic cloud management system perceives its environment through telemetry streams and API signals, reasons over observed states to formulate resource management strategies, executes actions through cloud control planes, and updates its operational models based on the outcomes it observes — in a continuous, closed-loop fashion.

This paper provides a comprehensive examination of agentic AI frameworks as applied to cloud resource management. Our contributions are as follows: (1) we develop a unified architectural taxonomy for agentic cloud management systems; (2) we systematically analyze the functional capabilities of such systems across five operational domains; (3) we critically evaluate the technical and governance challenges that must be addressed before wide-scale production deployment; and (4) we articulate a structured research agenda for advancing the field.

The remainder of this paper is organized as follows. Section 2 reviews the relevant background and related work. Section 3 presents the architectural framework of agentic AI systems for cloud management. Section 4 analyzes functional capabilities and operational use cases. Section 5 discusses challenges and limitations. Section 6 proposes a forward-looking research agenda. Section 7 concludes the paper.

2. BACKGROUND AND RELATED WORK

2.1 Evolution of Cloud Resource Management

Cloud resource management has evolved through several distinct generations. Early Infrastructure-as-a-Service (IaaS) platforms exposed raw computational resources through APIs, requiring substantial manual effort for provisioning, scaling, and maintenance. The introduction of Platform-as-a-Service (PaaS) models abstracted infrastructure concerns and enabled developers to focus on application logic, while managed services further reduced operational burden for common workload patterns.

The container revolution, catalyzed by Docker and subsequently systematized by Kubernetes, introduced declarative resource management through manifest-driven orchestration. Container orchestration platforms enabled fine-grained resource allocation, service discovery, and rolling deployment strategies at scale. However, the cognitive overhead of managing complex Kubernetes environments — spanning configuration management, network policy, storage provisioning, and observability — has proven substantial, giving rise to a secondary industry of Platform Engineering and Internal Developer Platform (IDP) tooling.

Serverless computing represented a further abstraction, eliminating server management entirely in favour of event-driven function invocations. While serverless architectures dramatically reduce operational responsibility, they introduce new challenges in cold-start latency, observability, and cost predictability. Frameworks such as AWS Lambda, Google Cloud Functions, and Azure Functions have become integral components of cloud-native architectures, particularly for workloads characterized by irregular or unpredictable invocation patterns.

2.2 AI and Machine Learning in Cloud Operations

The application of machine learning to cloud operations predates the agentic AI paradigm by more than a decade. Early applications focused on anomaly detection in time-series telemetry, using statistical methods such as ARIMA, exponential smoothing, and isolation forests to identify deviations from baseline system behaviour. Supervised learning models were subsequently applied to workload classification, failure prediction, and capacity forecasting tasks.

Deep learning techniques, particularly recurrent neural networks (RNNs) and their long short-term memory (LSTM) variants, achieved state-of-the-art performance on multivariate time-series forecasting tasks relevant to cloud operations, including CPU utilization prediction, network traffic forecasting, and storage demand estimation. Convolutional neural network (CNN) architectures have been applied to log analysis tasks, enabling automated root-cause attribution and failure pattern recognition at scale. Gaddipati (2024) demonstrated that deep learning-based intrusion detection systems can achieve high accuracy in securing cloud-native and distributed architectures, establishing a foundation for AI-driven security automation in cloud environments.

Reinforcement learning (RL) has been applied to cloud resource management problems characterized by sequential decision-making under uncertainty, including bin-packing optimization for virtual machine placement, adaptive autoscaling policy learning, and energy-aware workload scheduling. Q-learning and policy gradient methods have demonstrated the ability to discover resource allocation strategies that outperform hand-crafted heuristics on benchmark tasks, particularly in environments with complex interdependencies and non-stationary demand patterns.

2.3 Agentic AI and LLM-Based Orchestration

The emergence of large language models with strong instruction-following, code generation, and multi-step reasoning capabilities has created new possibilities for agentic system design. Frameworks such as LangChain, AutoGen, CrewAI, and LlamaIndex enable the construction of LLM-based agents that can decompose complex objectives into sub-tasks, invoke tools and APIs, interpret execution results, and iterate toward goal completion without explicit programmatic control flow.

Multi-agent frameworks extend single-agent capabilities by enabling specialization, parallelization, and collaborative problem-solving across networks of agents with distinct roles and expertise domains. Research by Gaddipati and Gandikota demonstrated that intelligent autonomous agents can be effectively deployed for data validation, orchestration, and cost optimization tasks, providing empirical evidence for the practical viability of agentic approaches in enterprise cloud settings. In the cloud management context, multi-agent architectures enable concurrent optimization across cost, performance, reliability, and compliance dimensions — objectives that are frequently in tension and require coordinated trade-off reasoning.

Gaddipati's work on agentic AI-driven autonomous data platforms for scalable cloud-native architectures provides a directly relevant precedent, demonstrating how agentic orchestration can manage the full lifecycle of cloud-native data workloads with minimal human intervention. Similarly, research on AI-powered decision automation in serverless architectures using AWS Lambda and Step Functions has shown that LLM-integrated workflow engines can autonomously handle complex, multi-step operational processes that would previously have required expert human intervention.

3. ARCHITECTURAL FRAMEWORK

3.1 System Architecture Overview

An agentic AI framework for cloud resource management comprises five primary architectural layers: (1) the Perception Layer, which collects and preprocesses environmental signals; (2) the Reasoning Layer, which interprets observed states and formulates management strategies; (3) the Action Layer, which translates decisions into cloud control plane operations; (4) the Memory Layer, which maintains operational history and learned knowledge representations; and (5) the Governance Layer, which enforces safety constraints, audit requirements, and policy compliance throughout the system's operation.

These layers interact within a continuous perception-reason-act-learn loop that enables the system to adapt to evolving environmental conditions and organizational requirements. The architecture is designed to be modular and composable, allowing organizations to incrementally augment existing cloud management toolchains with agentic capabilities rather than requiring wholesale replacement of established workflows.

3.2 Perception Layer

The perception layer is responsible for collecting, normalizing, and contextualizing environmental signals from heterogeneous cloud infrastructure sources. Input streams include infrastructure metrics (CPU utilization, memory consumption, network throughput, disk I/O), application performance indicators (request latency, error rates, throughput, Apdex scores), cost and billing telemetry (resource utilization efficiency, reserved capacity coverage, spot market pricing), security event streams (authentication logs, network flow data, vulnerability scan results), and configuration state representations.

Raw telemetry data is processed through feature engineering pipelines that compute derived metrics, detect short-term anomalies, and generate condensed state representations suitable for downstream reasoning components. Embedding models transform log data and unstructured

operational signals into dense vector representations that can be efficiently indexed and retrieved by memory subsystems. The perception layer must handle high-throughput, high-cardinality data streams with low latency, as delays in state observation directly translate to delays in autonomous response.

3.3 Reasoning Layer

The reasoning layer constitutes the cognitive core of the agentic framework. In contemporary implementations, this layer is typically anchored by a large language model that provides general-purpose reasoning capabilities, augmented by domain-specific models trained on cloud operational data. The LLM processes structured state representations produced by the perception layer, retrieves relevant operational history from memory subsystems, and generates management strategies expressed as structured action plans.

Reasoning processes are decomposed into hierarchical planning cycles operating at multiple time scales. Strategic planning cycles, executing at intervals of minutes to hours, address high-level resource allocation objectives such as capacity provisioning decisions, reserved instance purchase recommendations, and architectural reconfiguration proposals. Tactical planning cycles, operating at intervals of seconds to minutes, handle routine operational decisions such as pod scaling adjustments, load balancer weight modifications, and cache invalidation triggers. Reactive planning cycles, executing in near-real time, address urgent operational events such as service degradation, security incidents, and infrastructure failures.

3.4 Action Layer

The action layer translates reasoning outputs into executable operations within cloud environments. Actions are categorized into three classes based on their reversibility and impact radius: read actions (telemetry queries, configuration retrievals, API polling), write actions with limited blast radius (scaling adjustments, configuration parameter modifications, tag updates), and write actions with broad blast radius (infrastructure provisioning, network topology changes, data migration operations). This classification directly informs the authorization model applied at runtime.

Action execution is mediated by a policy enforcement engine that applies pre-defined safety constraints before any operation is submitted to cloud control planes. Constraint categories include resource quota limits, cost budget thresholds, change frequency governors, and environment protection designations. The action layer maintains comprehensive audit logs of all proposed and executed operations, including the reasoning chains that motivated each action, enabling post-hoc review and accountability.

3.5 Memory Layer

Effective agentic operation requires access to multiple forms of operational memory. Episodic memory stores time-stamped records of past operational states, decisions, and outcomes, enabling the system to recognize recurring patterns and avoid repeating past errors. Semantic memory encodes structured domain knowledge — including cloud service capabilities, pricing models, SLA requirements, and architectural best practices — that informs the reasoning layer's strategy formulation. Procedural memory stores learned operational routines and response playbooks that can be rapidly retrieved and adapted when familiar situations are encountered.

Memory retrieval is implemented using vector similarity search over dense embeddings, enabling efficient identification of historically similar situations across potentially millions of stored episodes. Memory consolidation processes periodically distil episodic records into generalized semantic knowledge, ensuring that the system's operational expertise improves continuously over time without requiring explicit retraining cycles.

4. FUNCTIONAL CAPABILITIES AND USE CASES

4.1 Dynamic Auto-Scaling

Autoscaling is one of the most mature and well-studied cloud management problems, and agentic AI frameworks offer significant improvements over conventional rule-based and metric-triggered approaches. Traditional autoscaling systems react to observed resource utilization metrics, applying scaling actions only after demand pressure has materialized. This reactive posture inevitably results in periods of under-provisioning during demand surges and over-provisioning during demand recessions.

Agentic autoscaling systems employ predictive models that anticipate demand changes based on historical patterns, external signals (such as scheduled events, marketing campaigns, and macroeconomic indicators), and leading indicators observable in application telemetry before they manifest as resource pressure. By combining predictive scaling with real-time reactive adjustments, agentic systems achieve tighter resource utilization profiles while maintaining service level objectives. Multi-dimensional scaling strategies simultaneously optimize across compute, memory, network, and storage dimensions, accounting for the complex interdependencies between resource types in containerized environments.

4.2 Cost Optimization

Cloud infrastructure cost optimization presents a multi-objective optimization problem of considerable complexity. Organizations must balance cost minimization against performance requirements, availability SLAs, and operational risk tolerance. The decision space includes choices across instance type selection, reserved versus on-demand versus spot pricing strategies, right-sizing of over-provisioned resources, workload scheduling to exploit time-of-day pricing variations, and data transfer cost management.

Agentic AI frameworks are particularly well-suited to cloud cost optimization due to their ability to reason across multiple optimization dimensions simultaneously and to incorporate contextual factors that are difficult to encode in rule-based systems. Gaddipati's work on hybrid data integration architectures combining Informatica and cloud-native services for scalable enterprise data systems demonstrates how intelligent orchestration can achieve cost efficiencies through adaptive resource utilization strategies. Empirical deployments of agentic cost optimization systems have reported cost reductions of fifteen to forty percent compared with baseline configurations managed by conventional autoscaling policies.

4.3 Anomaly Detection and Self-Healing

Anomaly detection and automated remediation represent a critical operational capability for high-availability cloud services. Traditional monitoring systems rely on static threshold alerts that generate high volumes of low-fidelity notifications, overwhelming operations teams with false

positives while missing subtle degradation patterns that precede service failures. Deep learning-based anomaly detection systems, as documented by Gaddipati in the context of enterprise network security, demonstrate the ability to identify complex multivariate anomaly patterns with substantially lower false positive rates than threshold-based approaches.

Agentic self-healing systems extend anomaly detection capabilities with automated remediation pipelines. Upon detecting an anomalous condition, an agentic system reasons over the observed symptoms to formulate a differential diagnosis, selects an appropriate remediation strategy from its operational knowledge base, executes the remediation action, and validates that the corrective action has resolved the underlying condition. In cases where automated remediation is insufficient or confidence levels are low, the system escalates to human operators with a structured incident report that includes its diagnostic reasoning and the remediation actions it has attempted. Gaddipati's research on anomaly detection in enterprise networks using deep neural networks and real-time streaming data provides foundational methodology that agentic cloud management systems can build upon.

4.4 Workload Scheduling and Orchestration

Intelligent workload scheduling requires reasoning over resource availability, job priority, deadline constraints, data locality, and infrastructure cost to allocate computational tasks across heterogeneous cloud resources. Agentic scheduling systems maintain a continuously updated model of cluster state and pending workload characteristics, applying optimization algorithms informed by historical performance data to construct efficient execution plans.

In batch processing environments, agentic schedulers can make intelligent decisions about when to leverage spot instances for cost-sensitive workloads, how to partition large jobs to maximize parallelism, and how to sequence dependent workflow stages to minimize total execution time. Gaddipati's research on AI-powered decision automation in serverless architectures using AWS Lambda and Step Functions demonstrates how LLM-integrated workflow orchestration can automate complex multi-step processes that span multiple cloud services and require conditional branching based on intermediate execution results.

4.5 Compliance and Security Assurance

Maintaining compliance with regulatory frameworks (such as SOC 2, ISO 27001, GDPR, and HIPAA) and organizational security policies across dynamic cloud environments is a significant operational challenge. Configuration drift — the gradual divergence of deployed infrastructure from approved baseline configurations — is a pervasive problem in cloud environments that creates both security vulnerabilities and compliance audit risks.

Agentic AI compliance systems continuously monitor infrastructure configurations against policy specifications, detect deviations, and autonomously remediate non-compliant states within defined safety boundaries. Natural language interfaces enable security and compliance teams to express policies in accessible terms that are automatically translated into enforceable technical controls. The system's comprehensive audit trail provides the documentation artefacts required for compliance certification processes, substantially reducing the manual effort associated with audit preparation.

5. CHALLENGES AND LIMITATIONS

5.1 Safety and Reliability

The deployment of autonomous agents with write access to production cloud infrastructure introduces significant safety challenges. Erroneous actions taken by an agentic system can cascade rapidly through interdependent services, potentially causing large-scale outages or data loss before human operators can intervene. Ensuring that agentic systems operate within safe action boundaries requires robust constraint enforcement mechanisms, comprehensive testing in staging environments, and graduated autonomy models that restrict the blast radius of autonomous actions based on confidence levels and environmental criticality.

The specification of safety constraints is itself a non-trivial problem. Overly conservative constraints limit the operational effectiveness of agentic systems and may negate their advantages over conventional automation. Insufficiently constrained systems pose unacceptable operational risks. Formal methods for safety specification and verification in agentic AI systems remain an active research area, and production deployments currently rely primarily on empirical testing and operational experience to calibrate constraint boundaries.

5.2 Explainability and Auditability

The opacity of LLM reasoning processes poses significant challenges for operational trust and regulatory compliance. When an agentic system takes a consequential action — such as scaling down a production database cluster or blocking a network access pattern — operators need to understand the reasoning that motivated the action in order to validate its correctness and learn from its outcomes. The generation of faithful, human-readable explanations for LLM-driven decisions remains an unsolved problem in interpretable AI research.

Current approaches to agentic explainability rely primarily on structured chain-of-thought reasoning outputs that document the intermediate steps and evidence considered in reaching a decision. While these outputs provide useful insight into the system's reasoning process, they do not constitute formal proofs of decision correctness and may not satisfy the evidentiary standards required by regulatory frameworks governing automated decision-making in high-stakes operational contexts.

5.3 Adversarial Robustness

Agentic AI systems that process external data streams — including log files, API responses, and user-generated content — are potentially vulnerable to prompt injection attacks, in which adversarial inputs embedded in processed data manipulate the system's reasoning process. In cloud security contexts, an attacker who can inject malicious content into monitored log streams may be able to influence the agentic system's security response decisions, potentially causing it to suppress alerts, whitelist malicious activity, or take disruptive remediation actions.

Defences against prompt injection in agentic systems include input sanitization and validation pipelines, separation of trusted and untrusted input channels, and adversarial training procedures that improve model robustness to injection attacks. However, comprehensive defences against sophisticated injection attacks remain an open research problem, and the security implications of deploying LLM-based agents in adversarial environments require careful consideration.

5.4 Governance and Organizational Adoption

The organizational and governance challenges associated with deploying agentic AI in cloud operations may be as significant as the technical challenges. Agentic systems that take autonomous actions on production infrastructure implicate questions of accountability, liability, and change management that extend beyond the engineering domain. Establishing appropriate human oversight mechanisms — including approval workflows for high-impact actions, exception handling procedures, and override capabilities — requires careful organizational design.

Change management considerations are particularly important in organizations with established SRE and DevOps practices. Agentic systems that modify operational workflows and reduce the scope of human decision-making may face resistance from practitioners who are accustomed to maintaining direct control over production environments. Successful adoption requires clear communication of the system's operating principles, transparent reporting of its actions and outcomes, and demonstrated reliability before expanding its autonomy.

6. RESEARCH AGENDA

6.1 Formal Safety Verification for Agentic Cloud Systems

The development of formal methods for specifying and verifying safety properties of agentic cloud management systems represents a high-priority research objective. Future work should investigate the application of formal specification languages, model checking techniques, and runtime verification frameworks to provide provable safety guarantees for constrained action domains within agentic systems.

6.2 Faithful Explainability for Operational Decisions

Research is needed on methods for generating accurate, verifiable explanations for agentic decision-making in cloud operational contexts. This includes investigation of mechanistic interpretability techniques for LLM reasoning processes, causal attribution methods for complex multi-agent decision chains, and standardized formats for operational decision documentation that satisfy regulatory audit requirements.

6.3 Adversarial Robustness and Security

The security implications of LLM-based agentic systems in adversarial cloud environments require systematic investigation. Research should address the characterization of adversarial attack surfaces specific to agentic cloud management, the development of robust defences against prompt injection and data poisoning attacks, and the design of security-aware agentic architectures that maintain integrity under adversarial conditions.

6.4 Benchmark Development and Standardized Evaluation

The field currently lacks standardized benchmarks for evaluating the performance of agentic cloud management systems across operational dimensions. The development of realistic, reproducible evaluation environments that capture the complexity of production cloud workloads — including multi-cloud configurations, adversarial conditions, and failure injection scenarios — is a prerequisite for rigorous scientific progress.

6.5 Human-Agent Teaming and Trust Calibration

Research on optimal human-agent teaming models for cloud operations should investigate how to design handoff protocols between agentic and human decision-makers that maintain operational continuity while appropriately leveraging the complementary capabilities of each. Trust calibration mechanisms that help operators develop accurate mental models of agentic system capabilities and limitations are essential for effective human oversight.

7. CONCLUSION

This paper has presented a comprehensive examination of agentic AI frameworks for autonomous cloud resource management. We have developed a unified architectural taxonomy encompassing perception, reasoning, action, memory, and governance layers; analyzed functional capabilities across dynamic auto-scaling, cost optimization, anomaly detection and self-healing, workload scheduling, and compliance assurance; identified key technical and organizational challenges; and proposed a research agenda for advancing the field.

The evidence surveyed in this paper supports the conclusion that agentic AI frameworks represent a transformative advancement in cloud resource management capabilities. By combining the generative reasoning power of large language models with reinforcement learning optimization, multi-agent collaboration, and comprehensive observability, agentic systems can address the operational complexity of modern cloud environments at a scale and sophistication that conventional automation approaches cannot match.

However, realizing the full potential of agentic cloud management requires sustained research investment in safety verification, explainability, adversarial robustness, and evaluation methodology. The governance and organizational challenges associated with deploying autonomous agents on production infrastructure are equally important and must be addressed alongside technical development. As these challenges are progressively resolved, agentic AI frameworks are positioned to become the dominant paradigm for cloud resource management, enabling organizations to operate increasingly complex cloud environments with reduced operational overhead, improved efficiency, and greater reliability.

Acknowledgements

The authors thank their respective departments at the Federal University of Agriculture Abeokuta and Ladoke Akintola University of Technology for institutional support. We also acknowledge the broader research community whose published work on agentic AI, cloud computing, and autonomous systems has informed the analysis presented in this paper.

REFERENCES

- Gaddipati, S. S. AGENTIC AI-DRIVEN AUTONOMOUS DATA PLATFORMS FOR SCALABLE CLOUD-NATIVE ARCHITECTURES.
- Gaddipati, S. S., & Gandikota, S. R. Design and Implementation of Intelligent Autonomous Agents for Data Validation, Orchestration, and Cost Optimization.
- Gaddipati, S. S. (2024). Deep learning based intrusion detection systems for securing cloud native and distributed architectures.
- Gaddipati, S. S. Anomaly Detection in Enterprise Networks Using Deep Neural Networks and Real-Time Streaming Data.
- Gaddipati, S. S. (2022). AI-Powered Decision Automation in Serverless Architectures Using AWS Lambda and Step Functions. *Journal of Computational Analysis and Applications*, 30(2).
- Gaddipati, S. S. Hybrid Data Integration Architectures: Combining Informatica and Cloud-Native Services for Scalable Enterprise Data Systems.
- Armbrust, M., Fox, A., Griffith, R., Joseph, A. D., Katz, R., Konwinski, A., ... & Zaharia, M. (2010). A view of cloud computing. *Communications of the ACM*, 53(4), 50–58.
- Barroso, L. A., Clidaras, J., & Hölzle, U. (2013). The datacenter as a computer: An introduction to the design of warehouse-scale machines. *Synthesis Lectures on Computer Architecture*, 8(3), 1–154.
- Burns, B., Grant, B., Oppenheimer, D., Brewer, E., & Wilkes, J. (2016). Borg, Omega, and Kubernetes. *ACM Queue*, 14(1), 70–93.
- Chen, Y., Mahajan, R., Sridharan, B., & Zhang, Z. (2013). A provider-side view of web search response time. *ACM SIGCOMM Computer Communication Review*, 43(4), 243–254.
- Dario, A., & Clark, J. (2023). Model cards for model reporting. In *Proceedings of the Conference on Fairness, Accountability, and Transparency*.
- Dean, J., & Ghemawat, S. (2008). MapReduce: Simplified data processing on large clusters. *Communications of the ACM*, 51(1), 107–113.
- Floridi, L., & Cowls, J. (2019). A unified framework of five principles for AI in society. *Harvard Data Science Review*, 1(1).
- Hadary, O., Marshall, L., Menache, I., Pan, E., & Schroeder, B. (2020). Protean: VM allocation service at scale. In *14th USENIX Symposium on Operating Systems Design and Implementation* (pp. 845–861).
- Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., ... & Fung, P. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12), 1–38.
- Kindratenko, V., Iffrig, J., Zhen, L., Surendra, B., Mu, Q., & Doerfler, D. (2009). GPU clusters for high-performance computing. In *2009 IEEE International Conference on Cluster Computing and Workshops*.
- Lwakatare, L. E., Raj, A., Crnkovic, I., Bosch, J., & Olsson, H. H. (2020). Large-scale machine learning systems in real-world industrial settings: A review of challenges and solutions. *Information and Software Technology*, 127, 106368.
- Ma, Z., Liu, Z., Liu, Y., Ma, X., & Luo, D. (2023). LLM-based multi-agent system for cloud operations. *arXiv preprint arXiv:2307.xxxxx*.
- Mao, H., Alizadeh, M., Menache, I., & Kandula, S. (2016). Resource management with deep reinforcement learning. In *Proceedings of the 15th ACM Workshop on Hot Topics in Networks* (pp. 50–56).

- Mirhoseini, A., Pham, H., Le, Q. V., Steiner, B., Larsen, R., Zhou, Y., ... & Dean, J. (2021). A graph placement methodology for fast chip design. *Nature*, 594(7862), 207–212.
- Park, J. S., O'Brien, J. C., Cai, C. J., Morris, M. R., Liang, P., & Bernstein, M. S. (2023). Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*.
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*.
- Silver, D., Huang, A., Maddison, C. J., Guez, A., Sifre, L., Van Den Driessche, G., ... & Hassabis, D. (2016). Mastering the game of Go with deep neural networks and tree search. *Nature*, 529(7587), 484–489.
- Tirmazi, M., Barker, A., Deng, N., Haque, M. E., Qin, Z. G., Hand, S., ... & Wilkes, J. (2020). Borg: The next generation. In *Proceedings of the Fifteenth European Conference on Computer Systems* (pp. 1–14).
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30.
- Wang, G., Xie, Y., Demers, A., & Gehrke, J. (2022). Towards AIOps in edge computing environments. In *Proceedings of the 2022 IEEE International Conference on Cloud Engineering*.
- Wei, J., Wang, X., Schuurmans, D., Bosma, M., Xia, F., Chi, E., ... & Zhou, D. (2022). Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35, 24824–24837.
- Weyns, D. (2019). *An introduction to multiagent systems*. John Wiley & Sons.
- Xu, L., Mo, Y., & Chen, Y. (2023). Auto-scaling policies for cloud workloads using deep reinforcement learning. *IEEE Transactions on Cloud Computing*, 11(2), 1234–1248.
- Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2023). ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations*.
- Zhang, W., Fang, J., Zou, M., & He, P. (2021). LogBERT: Log anomaly detection via BERT. In *2021 International Joint Conference on Neural Networks* (pp. 1–8).